

Natural Language Processing

Natural language processing, or NLP, focuses on enabling computers to process and understand human language.

Tokenization divides text into smaller units called tokens.

Text classification assigns labels to text. Sentiment analysis is an example of text classification.

Word embeddings represent words or tokens as numerical vectors that capture useful semantic relationships.

Transformers use attention mechanisms to model relationships between tokens in a sequence.

Large language models are neural language models trained on large text datasets to perform tasks such as generation and question answering.

Retrieval can provide relevant external documents to an NLP system before an answer is generated.