

Retrieval-Augmented Generation

Retrieval-augmented generation, or RAG, combines information retrieval with language generation.

A RAG system first retrieves relevant documents from a knowledge source and then provides those documents as context to a language model.

BM25 is a sparse retrieval method based primarily on term matching and term frequency statistics.

Dense retrieval represents queries and documents as vectors and compares their semantic similarity.

Embeddings are numerical vector representations of text that can be used for semantic retrieval.

Chunking divides long documents into smaller passages before indexing.

Good retrieval quality is important because the generator can only use the retrieved context that is supplied to it.